

$L = RL = BPL$

An expository account of the proof

A rewrite, for human readers, of the manuscript
Exact Derandomization of Logarithmic Space: $L = RL = BPL$
(OpenAI, September 23, 2026)

Prepared with Claude (Anthropic) · October 7, 2026

Abstract

In September 2026 OpenAI posted a 108-page manuscript, written by an internal model, claiming a deterministic logarithmic-space simulation of every bounded-error randomized logarithmic-space computation that runs in polynomial time. In fact it approximates the acceptance probability itself to any fixed inverse-polynomial accuracy in deterministic logarithmic space. The manuscript is complete in the sense that every step is written down, but it is very hard to read: ideas are introduced exactly when needed and rarely motivated. This note explains the architecture of the proof for a complexity theorist. The first, purely algebraic, half (a “correction and copy hierarchy” that pushes all acceptance probability into a reward vector in $O(\log n)$ stages) is elementary and is presented with full proofs. The second half (estimating that reward with $O(\log n)$ shared random bits, then enumerating them, all in $O(\log n)$ space) is a long chain of interlocking devices; for each one this note explains what problem it solves, states what is proved, and points to the manuscript. Unlike the Unique Games result in the same release, this one has no Lean formalization. [Section 6](#) says what was checked here and where a referee should concentrate.

Contents

1	Introduction	3
1.1	The claims	3
1.2	Context	3
1.3	Why the problem is hard, and the two halves of the answer	3
1.4	The proof on one page	4
1.5	Status	5
2	The exact target: the correction and copy hierarchy	5
2.1	Acceptance probability as a matrix inverse	5
2.2	Moving mass into the reward	6
2.3	Overloaded columns and detectors	7
2.4	Copies	7
2.5	The target	8
3	Estimating the target with a shared random environment	8
3.1	What has to be estimated	9
3.2	Atoms: sparse tables with exact conditional means	9

3.3	Conditional averaging with property (T)	10
3.4	The budget schedule	10
3.5	Compression of rows	11
3.6	Fingerprints within budget blocks	11
4	Fitting the estimator into logarithmic space	12
4.1	What is kept, and what is shared	12
4.2	The potential argument	13
5	Finishing the proof	14
5.1	From an average to one start, and the median	14
5.2	Consequences	14
5.3	The order of constants	15
6	A guide for the careful reader	15
6.1	What this note does and does not establish	15
6.2	Where to look hardest	15
6.3	Correspondence with the manuscript	16

1 Introduction

1.1 The claims

Machines have a read-only input, finitely many work tapes, and fresh fair coins; space counts every work cell visited. RL and BPL are the languages decided by polynomial-time randomized $O(\log n)$ -space machines with one-sided error (acceptance probability 0 versus $\geq 1/2$) and two-sided error ($\leq 1/3$ versus $\geq 2/3$). Write $B = \Theta(\log(n+2))$ throughout.

Theorem 1.1 (Theorem 1.1 of [1]). $L = RL = BPL$.

The proof does not use the gap. It approximates the acceptance probability.

Theorem 1.2 (Theorems 1.2 and 12.1 of [1]). *Let $\mathcal{M}$ be a randomized machine running in polynomial time and $O(\log n)$ space, with acceptance probability $p_{\mathcal{M}}(x)$. For every fixed d there is a deterministic $O(\log n)$ -space, polynomial-time transducer that outputs a dyadic z_x with $|z_x - p_{\mathcal{M}}(x)| \leq (n+2)^{-d}$. With the accuracy given in unary as 1^q , one algorithm achieves error 2^{-q} in space $O(\log n + q)$ and time $n^{O(1)}2^{O(q)}$.*

Consequences proved in the manuscript include $\text{PromiseBPL} = \text{PromiseL}$, deterministic log-space construction of an accepting computation whenever the acceptance probability is at least inverse-polynomial (Section 5.2), and an effective compiler that turns a randomized machine with stated resource bounds into a deterministic one with explicit bounds.

1.2 Context

The best previous general results were Nisan's $O(\log^2 n)$ -space pseudorandom generator [9], his simultaneous polynomial-time, $O(\log^2 n)$ -space simulation [10], the Saks–Zhou $O(\log^{3/2} n)$ -space simulation [12], and Hoza's improvement to $O(\log^{3/2} n / \sqrt{\log \log n})$ [7]. Reingold's theorem [11] settled undirected connectivity, and Cheng and Hoza [5] showed that log-space hitting sets for read-once branching programs would give $L = BPL$. The present proof does not build a pseudorandom generator or a hitting set. It *computes* the acceptance probability, using randomness only inside an estimator whose $O(\log n)$ random bits are finally enumerated.

1.3 Why the problem is hard, and the two halves of the answer

The acceptance probability is a coordinate of $(I - S)^{-1}e$, where S is the transition matrix of the configuration graph. Computing it means summing over polynomially long random paths. Naive recursion (Savitch-style squaring) uses $O(\log n)$ levels of $O(\log n)$ bits each, hence $O(\log^2 n)$ space. Every improvement has to make the levels cost less.

The manuscript separates two problems.

1. **An exact algebraic target with $O(\log n)$ levels** (Section 2). A deterministic sequence of transformations (C_l, W_l) of the transition matrix and a reward vector, the *correction and copy hierarchy*, moves probability mass out of the transitions and into the reward. Each stage divides every transition entry by 2^H while at most doubling the number of active vertices. After $L = O(\log n)$ stages the transitions carry negligible mass, and the reward at one designated vertex equals the acceptance probability up to $2^{-\Omega(L)}$. This part is short,

elementary, and fully checkable; its statements involve expectations over random hash functions, but those are deterministic numbers.

2. **Evaluating the target in $O(\log n)$ space** (Sections 3 and 4). The hierarchy’s objects are polynomial-size matrices that cannot be stored, and each stage depends on the previous one. The manuscript builds a randomized estimator of the final reward that (a) uses a single shared random “environment” of $O(\log n)$ bits for all levels at once, (b) is accurate for most environments, and (c) can be evaluated deterministically, for each fixed environment, in $O(\log n)$ space. Enumerating all $2^{O(\log n)}$ environments and taking a median finishes the proof (Section 5).

The key accounting idea. The recursion in part 2 is not organized by hierarchy stage but by a *statistical budget* $M = O(\log n)$. A task at budget M is computed as a telescoping sum of differences between child estimates at budgets m and $m - 1$, and a child at “slack” $d = M - m$ may retain $O(d)$ bits while it runs. Since the budget drops by at least d when $O(d)$ bits are spent, the bits retained along any chain of nested calls telescope to $O(M_{\text{top}}) = O(\log n)$. Everything else in part 2 (sparse sampled tables, expander-walk averaging, compression, fingerprints, a catalytic bit vector, tree-contour traversal, digit-by-digit arithmetic) exists to make this one accounting rule true.

1.4 The proof on one page

1. **Configuration matrix** (Section 2.1). Time-stamped configurations give a strictly forward substochastic S with bounded in- and out-degree, and $p_0 = (I - S)^{-1}e$ holds the acceptance probability from every configuration.
2. **Correction** (Section 2.2). For a nonnegative E , setting $D = C - E + EC$ gives $I - D = (I + E)(I - C)$, so moving E out of the transition and into the reward ($W \mapsto W + EW$) is exact. E is chosen to remove the large part of each large entry, except in “overloaded” columns with many large entries.
3. **Copies** (Section 2.4). Overloaded columns are split into D_0 copies, each receiving $1/D_0$ of the large entries. Detectors defined by random sampling certify that overloaded columns are few, so active vertices at most double while entries shrink by 2^H .
4. **Target** (Section 2.5). After $L = O(\log n)$ stages, $0 \leq p_0(x) - W_L(\bar{x}_L) \leq |V_0|2^{-(H-1)L}$.
5. **Environment and atoms** (Sections 3.1 and 3.2). Each matrix entry is represented by a sparse random table on retained vertex sets chosen by polynomial hashing over a field of size $2^{O(\log n)}$. Its conditional mean is exactly the entry. Errors are measured in an L^8 norm integrated over sources and environments.
6. **Averaging** (Section 3.3). Conditional means are approximated by walks on the environment space. An action of $\text{SL}_H(\mathbb{Z}) \ltimes \mathbb{Z}^H$, which has property (T), gives a spectral gap uniform in the field size.
7. **Budgets** (Section 3.4). Additive schedule, compression of rows to short lists, and short fingerprints of endpoints within budget blocks (Sections 3.5 and 3.6) prove an integrated error bound $C2^{-p}$ for the actual estimator (Theorem 3.3).
8. **Space** (Section 4). One cursor vertex, paired ports, tree contours for incoming edges, a




Figure 1: Architecture. The left column is Section 2; the right column is Sections 3 and 4.

shared catalytic bit vector for key equality, and redundant signed-digit arithmetic implement every query in $O(\log n)$ space on every environment.

9. **Finish** (Section 5). Pay a factor $N = 2^{O(\log n)}$ to pass from the integrated bound to the start vertex, enumerate the environments conditioned on retaining it, and take a median.

1.5 Status

The manuscript is 108 pages. Its first half is the hierarchy and the statistical estimator; its second half is the implementation (numerical streams and the controller), the conclusion and the compiler. The repository’s index lists no Lean formalization for this result, so, unlike the Unique Games proof in the same release, it rests on the written argument alone. The README notes that unformalized results “could have issues.” Given the claim, this paper deserves an expert referee; Section 6 suggests where to start.

2 The exact target: the correction and copy hierarchy

2.1 Acceptance probability as a matrix inverse

Lemma 2.1 (Configuration reduction; Lemma 2.1 of [1]). *For a polynomial-time randomized log-space machine and input x one can specify in log space a polynomial-size vertex set V_0 (time-stamped configurations), a strictly forward substochastic matrix S (every nonzero entry goes from time t to time $t + 1$; rows sum to at most 1) with constant in- and out-degree and paired ports, an indicator e of accepting configurations at the final time, and a start x_0 , such that $p_0 = (I - S)^{-1}e = \sum_j S^j e$ lies in $[0, 1]^{V_0}$ and $p_0(x_0)$ is the acceptance probability.*

Bounded in-degree (predecessors are found by guessing a constant amount of local change) and *paired ports* (from a vertex and an out-port one can compute the neighbor and the corresponding in-port, and conversely) will matter later: they let the algorithm move a single cursor along an edge and back without storing the old address.

2.2 Moving mass into the reward

The whole hierarchy rests on one identity. For square matrices C and E ,

$$D := C - E + EC \quad \implies \quad I - D = (I + E)(I - C).$$

So if $(I - C)p = W$, then $(I - D)p = (I + E)W$: the same vector p solves the new system with transitions D and reward $W + EW$.

Reading the identity as path shortcutting. Think of $E(x, y)$ as the part of the transition $x \rightarrow y$ that we decide to “pay out now”. The new reward at x gains $E(x, y)W(y)$, the reward collected at y ; and the new transitions gain $E(x, y)C(y, \cdot)$, which continue directly from y ’s successors. Removing E from C and adding EC is exactly replacing the edge $x \rightarrow y$ by “collect y ’s reward and jump past y ”. Nothing is approximated; the identity is exact for every E .

The point of a stage is to choose E so that D has *small entries*. Fix a stage with vertex set V , a rate $f = 2^{-Hl}$, and an input C that is strictly forward, substochastic, and has entries at most f . Write $q = 2^H$ ($H \geq 4$ fixed). A large entry of C can be removed only if E absorbs most of it, but E also creates new entries EC ; those can again be large. The manuscript solves the self-consistent equation

$$E_*(x, y) = f Q((C + E_*C)(x, y) / f) g(y),$$

where Q is a smooth cutoff with $(t - \theta_2)_+ \leq Q(t) \leq (t - \theta_1)_+$ for small fixed $0 < \theta_1 < \theta_2 < 1/(32q)$, and $g: V \rightarrow [0, 1]$ is a column multiplier (which columns are allowed to be corrected). With $g \equiv 1$, the residual $T = C - E_* + E_*C$ has every entry at most $\theta_2 f$.

Lemma 2.2 (Factorial truncation; Lemma 4.1 of [1]). *The fixed point E_* exists and is unique (define each row in increasing column order). The iterates $E_0 = 0$, $E_i = fQ((C + E_{i-1}C)/f)g$ increase to E_* , and*

$$\max_{x,y} \frac{E_*(x, y) - E_k(x, y)}{f} \leq e^{1/\theta_1} \frac{(L_Q/\theta_1)^k}{k!}.$$

So a fixed number k_* of iterations, independent of $|V|$, gets within $f/(32q)$ of the fixed point.

Why. The residual has rows summing to at most 1 and every positive entry of E_* leaves a residual of at least $\theta_1 f$, so row x of E_* has at most $m_x \leq 1/(\theta_1 f)$ nonzero entries. Unrolling $E_* \leq C + E_*C$ along row x writes $E_*(x, y)$ as a sum over increasing chains $x < z_1 < \dots < z_j < y$ inside that support set, each a product of $j + 1$ entries of size $\leq f$. There are $\binom{m_x}{j}$ chains, so $E_*(x, y) \leq f(1 + f)^{m_x} \leq f e^{1/\theta_1}$. The same unrolling for $E_* - E_k$ uses only chains of length exactly k , giving $\binom{m_x}{k} f^k \leq (1/\theta_1)^k / k!$. $\square$

2.3 Overloaded columns and detectors

Correction is not allowed everywhere. If many rows have large entries into the same column y , correcting all of them would later require handling all those incoming entries at y , and the estimator needs bounded incoming degree. So the multiplier $g(y)$ is turned off at *overloaded* columns, and these are found by a detector.

Sampling channels. Vertices have numerical IDs of $O(\log n)$ bits. A *rank channel* is a random $H \times h_0$ matrix $A = (a_0, F)$ over a prime field $\mathbb{F}_P$, $P = 2^{\Theta(\log n)}$, with F of full column rank. A vertex y is *retained at rate* $b_k = 2^{-Hk}$ if all H coordinates of $A(1, y, \dots, y^{h_0-1})^\top$ lie in $\{0, \dots, \lceil P2^{-k} \rceil - 1\}$. This is a pairwise-independent-style sample with

$$\Pr[\text{retained}] = \pi_b \asymp b, \quad \Pr[y \text{ retained} \mid x \text{ retained}] \leq C b \ (x \neq y),$$

and Chebyshev tails for counts (Lemma 3.1 of [1]). Channels X (sources), Y (terminals) and a few more are independent. The whole collection, plus some auxiliary bits, is the *environment* σ ; it has $O(\log n)$ bits.

The detector. At pilot step i let $G_i = C + \widehat{E}_{i-1}C$. A row gate $R_i(x)$ is 1 when row x has fewer than K retained large entries. The column score

$$U_i(y) = \mathbb{E} \left[\text{rank}_J \{ \phi(G_i(x, y)/f) R_i(x) : X_f(x) = 1 \} \mid Y_f(y) = 1 \right]$$

is a *deterministic* number: an expectation over the channels of the J -th largest score among sampled rows. Set $\widehat{g}_i(y) = \gamma(\max_{r \leq i} U_r(y))$, where the cutoff γ is 1 for scores up to $1/2$ and 0 from $4/5$ on, and after k_* pilot steps use $g = \widehat{g}_{k_*}$ in the actual iteration.

Lemma 2.3 (Detectors; Lemma 4.3 of [1]). *If more than R_*/f rows have $G_i(x, y) \geq \delta_1 f$, then $U_i(y) \geq 19/20$. If $U_i(y) > 1/4$, then more than $J/(4C_\pi f)$ rows have $G_i(x, y) > \delta_0 f$. Hence a corrected column ($g(y) > 0$) has at most R_*/f large incoming entries.*

The first statement is sampling: about $R_*\pi_f/f \gg J$ heavy rows are retained, and most have full row gates. The second is Markov's inequality. Thresholds are smooth cutoffs so that everything is later Lipschitz in its inputs.

Lemma 2.4 (Residual; Lemma 4.4 of [1]). *With $E = E_{k_*}$ and $D = C - E + EC$: D is nonnegative, strictly forward and substochastic; $D(x, y) \leq (k_* + 1)f$; and $D(x, y) \leq \varepsilon_{\text{good}} f < f/(16q)$ in every column with $g(y) = 1$. Moreover $D \geq c(C + E + EC)$ for a fixed $c > 0$.*

The last inequality says the three unsigned terms of D are each dominated by D , so later each can be estimated separately with error measured relative to D .

2.4 Copies

Entries of D are now small in good columns but can still be about f in overloaded ones. Make $D_0 > K_*q$ copies $y_0, \dots, y_{D_0-1}$ of every vertex, and route large entries evenly:

$$C^+(x_i, y_j) = r_i(x) D(x, y) s_j(D(x, y)/f),$$

where $s_0 \approx 1$ and $s_j \approx 0$ for small entries, $s_j = 1/D_0$ for all j when the entry exceeds $f/(4q)$, and the row factor $r_j(x)$ ($j > 0$) is 1 only for vertices whose detector score was high. Let $P^r(x_i, x) = r_i(x)$.

Proposition 2.5 (One stage; Prop. 4.5 of [1]). C^+ is strictly forward and substochastic, $C^+ \leq f/q$ entrywise, $C^+ P^r = P^r D$, and $|\mathcal{A}(C^+)| \leq 2|\mathcal{A}(C)|$, where $\mathcal{A}$ is the set of vertices touching a positive entry.

Why. Entries: a small entry stays below $f/(4q)$; a large one is split D_0 ways and is at most $K_* f/D_0 < f/q$. Intertwining: on any positive entry the routing fractions times row factors sum to 1, because a large entry in a column forces that column to be uncorrected, hence high-scoring, hence $r_j = 1$. Active count: only high-scoring parents $\mathcal{O}$ get extra active copies. Each has more than $J/(4C_\pi f)$ incoming entries above $\delta_0 f$, while the total number of such entries is at most $|\mathcal{A}(C)|k_*(k_* + 1)/(2\delta_0 f)$ (row sums are at most k_*). With J large, $(D_0 - 1)|\mathcal{O}| \leq |\mathcal{A}(C)|$. $\square$

2.5 The target

Start from $C_0 = S$, $W_0 = e$, and iterate with $f_l = q^{-l}$, $V_{l+1} = V_l \times \{0, \dots, D_0 - 1\}$, $C_{l+1} = C_l^+$, and

$$W_{l+1}(x_i) = r_i(x)(W_l(x) + (E_l W_l)(x)).$$

Proposition 2.6 (Prop. 4.7 of [1]). With P_l^c the cumulative projection, $W_l = P_l^c p_0 - C_l P_l^c p_0$ and $0 \leq W_l \leq 1$. For the all-primary copy $\bar{x}_l$ of a base vertex x ,

$$0 \leq p_0(x) - W_l(\bar{x}_l) \leq |V_0| 2^{-(H-1)l}.$$

Hence $L = O(\log n)$ stages give any fixed inverse-polynomial accuracy, and every V_l has $2^{O(\log n)}$ vertices with $O(\log n)$ -bit addresses.

Proof. By induction $C_l P_l^c = P_l^c - B_l(I - S)$ with $B_{l+1} = P_l^r(I + E_l)B_l$, using $C^+ P^r = P^r D$ and $I - D = (I + E)(I - C)$. Apply this to p_0 and use $(I - S)p_0 = e$. Every row of C_l has at most $2^l |V_0|$ positive entries, each at most q^{-l} , so its sum is at most $|V_0| 2^{-(H-1)l}$, and $0 \leq P_l^c p_0 \leq 1$ with $(P_l^c p_0)(\bar{x}_l) = p_0(x)$. $\square$

What the hierarchy buys. Every stage is exact, and the number of stages is $O(\log n)$. The probability does not have to be computed by multiplying out long paths; it has been “pushed” into a reward vector by a logarithmic number of local, sparse corrections. The price is that each stage is defined through the previous one, through nonlinear cutoffs, and through expectations over hash functions. The rest of the paper is about evaluating this object without $O(\log n)$ bits per stage.

3 Estimating the target with a shared random environment

This section and the next describe how $W_L(\bar{x}_L)$ is evaluated. Here the question is statistical: what is computed, and why is it accurate for most environments? [Section 4](#) explains why each computation fits in $O(\log n)$ space. Proofs in the manuscript are long; this section aims to make the statements and their purposes clear, with the main estimates indicated.

3.1 What has to be estimated

At each stage the hierarchy uses a fixed list of *tasks*, in a fixed dependency order: entries of C_l ; at each pilot step the detector scores and the vector U_i ; the pilot and actual corrections E_i ; entries of D_l routed to copies; and the reward update W_{l+1} . A *row task* is a whole sparse table (an estimate of one of these matrices); a *vector task* is a vertex function. The number of task types per stage is a constant, so there are $O(L) = O(\log n)$ task ranks in all.

All tasks share one random *environment* σ : a constant number of independent rank channels (Section 2.3) over a field of size $P = 2^{\Theta(\log n)}$, plus $O(\log n)$ auxiliary “salt” bits used for fingerprints (Section 3.6). The environment is shared by every task at every level. No independence between the errors of different tasks is ever assumed; all bounds are proved by triangle inequalities, Hölder, and caps on degrees.

3.2 Atoms: sparse tables with exact conditional means

A matrix entry $G(x, y)$ at stage rate f is represented on the slice where the source is retained at rate u and the terminal at rate a (written $X_u(x) = Y_a(y) = 1$). Entry sizes are split into scale bins $G = \sum_k G \omega_k(G)$, with ω_k a smooth partition of unity on dyadic scales $b_k = 2^{-Hk}$. At scale a the *dense atom* is

$$\mathcal{B}_a(x, y) = \mathbf{1}_{\{Y_t(y)=1\}} \frac{k_a(G(x, y)/a) R(x)}{c(q_y)},$$

where k_a is the bin function, $R(x) \in [0, 1]$ is a row gate that vanishes when row x has K or more retained large entries, and q_y is the exact conditional mean of the gate (a “preparatory” table), clipped to $[1/2, 1]$.

Lemma 3.1 (Atom identities; Lemmas 5.1 and 5.3 of [1]). *On the slice $X_u(x) = Y_t(y) = 1$, the conditional mean of a $\mathcal{B}_a(x, y)$ is exactly $G(x, y)\omega_k(G(x, y))$. Each row of an atom table has fewer than K nonzero entries. Correction atoms carry an additional incoming gate, so each column of an exposed correction relation has at most K_c entries; their conditional means are exactly the corrections.*

Products are estimated by sampling midpoints: an entry of EC is represented by $\frac{f}{\pi_f} \sum_{z:Z_f(z)=1} \mathcal{F}(x, z) \mathcal{B}_b(z, y)$, which has the right conditional mean because the channels X, Z, Y are independent and the true atoms factor through them (Lemma 7.1 of [1]).

Why sparse tables, and why exact means. Bounded row and column support is what lets a single cursor walk along table entries (Section 4). Exact conditional means are what make the estimator’s errors telescope in Section 3.4: the target of every sample term is a deterministic number, so the error can be split into a deterministic bias and a mean-zero noise, which are bounded by different arguments.

The error metric. For a row quantity Z and $s = 8$,

$$\|Z\|_{u,a;s}^s = \frac{1}{N} \mathbb{E}_\sigma \sum_x \frac{X_u(x)}{u} \sum_{y:Y_a(y)=1} |Z(x, y)|^s, \quad N = 2^{C_N \log n} \text{ (a bound on stage size),}$$

and errors of bounded nonnegative outputs are measured in the root discrepancy $\|H^{1/s} - (H')^{1/s}\|$. Three features matter: the error is *averaged over sources* (weight X_u/u , normalized by $1/N$); it is an L^8 norm over the environment; and taking 8th roots turns multiplicative errors into additive ones. At the very end, passing from an average over sources to one specific start costs a factor $N = n^{O(1)}$ (Section 5), which is paid by making the statistical precision a large constant times $\log n$.

The cost of releasing a source. In a product $\mathcal{F}(x, z)\mathcal{B}_b(z, y)$ the right factor's error is measured with sources z at rate f , not x at rate $u \leq f$. Because correction atoms have at most K_c incoming entries, each z is reached from at most K_c sources, and the change of source domain costs a factor f/u in the s -th power of the norm (Lemma 7.3 of [1]). This is paid by the precision convention below.

3.3 Conditional averaging with property (T)

Every atom identity is a conditional mean on a slice of environments. To approximate such a mean for a function $F(\sigma)$, the estimator replaces $F(\sigma)$ by the average of F over the endpoints of all walks of length d from σ (a deterministic sum over $2^{O(d)}$ walk words). If the walk has a spectral gap on the slice, this is close to the conditional mean.

Lemma 3.2 (Uniform gap; Lemmas 6.1–6.2 of [1]). *Let $G_H = \text{SL}_H(\mathbb{Z}) \ltimes \mathbb{Z}^H$ with $H \geq 3$ and a fixed finite symmetric generating set. The lazy walk has a spectral gap, depending only on H , on every finite transitive G_H -set. Consequently there is a Markov operator K on each conditional slice of environments with $\|K^d - \Pi\|_{L_s \rightarrow L_s} \leq C_s 2^{-c_s d}$, uniformly in the field size, the box sizes, and the conditioning vertices.*

The action is on a channel's coordinates (F, w) , where F is the frame and $w = A(1, y, \dots, y^{h_0-1})^\top$ is the retained box coordinate of the conditioning vertex y : (M, c) acts by $(F, w) \mapsto (MF, Mw + c)$. It is transitive, so the gap follows from property (T) of $\text{SL}_H(\mathbb{Z})$ and relative property (T) of $(\text{SL}_2(\mathbb{Z}) \ltimes \mathbb{Z}^2, \mathbb{Z}^2)$, with Shalom's explicit Kazhdan constants [13]; Riesz–Thorin moves the L_2 gap to L_8 .

Why property (T) instead of an ordinary expander. The algorithm must average over environments conditioned on retaining vertices it does not currently remember. So each walk step must be a fixed, measure-preserving transformation of the environment that does not mention the conditioning vertex, with validity checked only when the vertex is visited (Lemma 6.3: “guessed wraps” are enumerated and exactly one survives). A group action does this naturally, and property (T) gives one spectral gap for all the finite quotients that arise, whatever the field size and box sizes.

3.4 The budget schedule

Each task has an integer *statistical budget* M and an implicit precision

$$p = \frac{M}{K_{\text{prec}}} - \frac{1}{2} \log_2 \frac{1}{a} - \frac{1}{2} \log_2 \frac{1}{u} - A_* i_*,$$

where a, u are its end and start rates and i_* its rank in the dependency order. Tasks with $p \leq 0$ return zero. The inductive claim is that every task's error is at most $C_b 2^{-p}$. The $\frac{1}{2} \log$ terms are exactly what pays for scale changes (a bin at scale b costs $(b/a)^{1/8}$ or $(b/a)^{7/8}$, and weight $1/2$ sits between these) and for releasing a source (cost $(f/u)^{1/8}$).

Additive estimation. A task at budget M is computed as

$$\text{exact base terms} + \sum_b \sum_{m=1}^{M-c_b} K^{\ell(M-m)} (T_{b,m} - T_{b,m-1}),$$

where $T_{b,m}$ is a sample term using child estimates at budget m , $c_b \approx c_0(1 + |\log_2(b/a)|)$ cuts off scales far from a , and $K^{\ell(M-m)}$ is the conditional walk of length proportional to the slack $M - m$. The conditional means telescope to the term at the largest available budget, whose bias is bounded by induction (Lemma 7.6). The noise of each difference is small because both terms approximate the same quantity (Lemma 7.7 gives $2^{-p+C_0(M-m)}$ before averaging), and the walk contracts it by $2^{-c\ell(M-m)}$, which wins for a large fixed ℓ (Lemma 7.8).

A multilevel estimator. This is in the spirit of multilevel Monte Carlo: accurate children are used sparingly and with little averaging, cheap children often and with much averaging. Its real purpose here is space. A term at slack $d = M - m$ is described by $O(d)$ bits (walk word, scale offset, formula choices; Lemma 7.5), and calls a child of budget at most $M - d$. That is the exchange rate the space accounting of [Section 4](#) needs.

3.5 Compression of rows

Adding many increments can make a row array dense. After each dyadic slack group d_j , the array is replaced by a common default plus a list of at most $n_j = 2^{c_2 d_j}$ exceptions: the threshold is the largest of the old default, a floor $\rho_j = 2^{-c_1 d_j}$, and the n_j -th largest entry size; everything below twice the threshold becomes the default (Lemma 8.1). The map is uniformly Lipschitz.

The analysis (Section 8) separates three cases. Before the first “overflow” (threshold above the floor), all changes are tiny. An overflow caused by genuinely many heavy endpoints is rare: total row mass is bounded, so the expected number of retained heavy endpoints, even conditioned on one endpoint, is $O(1/\rho_j)$, and Markov gives probability $O(1/(n_j \rho_j))$. An overflow caused by errors at light endpoints is charged to those errors. With c_2 large relative to c_1 , all the resulting series converge (Prop. 8.5), so compression costs only a constant factor in the error. Range clamps (Lemma 8.6) keep numbers in $O(d)$ bits without changing any relevant output.

3.6 Fingerprints within budget blocks

Comparing two endpoint IDs exactly needs $O(\log n)$ bits, but a frame at small slack d may keep only $O(d)$ bits. For slack $d \leq 2h_*$, where $h_* = \lfloor \varepsilon \log_2 B \rfloor = \Theta(\varepsilon \log \log n)$, endpoints are therefore compared by keys $k(v) \in \{1, \dots, B\}$: $k(v) = \tau(\alpha P_v(r) + \beta)$ over a prime near B^3 , with P_v the polynomial whose coefficients are the bits of v . Distinct fixed IDs collide with probability $O(1/B)$ (Lemma 9.1).

The difficulty is *adaptivity*: which endpoints get compared depends on earlier computations, which may depend on the key. The manuscript handles this as follows (Section 9).

- Budgets are cut into blocks of h_* consecutive values, each with an independent salt. A task in block b may look only at salts of blocks $\leq b$; its walks vary the earlier salts and freeze the rest.
- A *clean calculation* uses exact ID comparison inside the current block. It is independent of the block’s salt (Lemma 9.2).
- The clean calculation determines an *envelope* of at most $2^{O(h_*)}$ candidate endpoints per row, independent of the salt, so all keys in it are distinct except with probability $2^{O(h_*)}B^{-1/2}$ (Lemma 9.3).
- Actual candidates outside the envelope are charged to child errors, using smooth “witness” functions (Lemma 9.4); collisions inside it are charged directly (Lemma 9.6); bucket sums only alias increments, never stored values (Lemma 9.5).
- A coupled induction bounds the actual–clean discrepancy by $2^{-p}2^{-D_1h_*+D_2(M-M_b)}$, closed by choosing ε small (Prop. 9.7).

Theorem 3.3 (Estimator accuracy; Theorem 9.8 of [1]). *Every task output satisfies the error bound $C_b 2^{-p}$. In particular,*

$$\frac{1}{N} \mathbb{E}_\sigma \sum_{x \in V_l} \frac{X_u(x)}{u} |\bar{W}_l(x) - W_l(x)|^8 \leq C_b^8 2^{-8p},$$

with constants independent of the input and of $\log n$.

4 Fitting the estimator into logarithmic space

Fix an environment σ . Every estimator value is now a deterministic finite recursion. The manuscript’s Sections 10–11 show that a single deterministic machine evaluates any requested binary digit of any task value in $O(\log n)$ space, on *every* environment, including the ones where the estimate is inaccurate. Totality matters: the final enumeration must finish on every environment.

4.1 What is kept, and what is shared

The shared facilities, present once for the whole computation, are: the input; the master environment ($O(\log n)$ bits); a single *cursor* holding one vertex ID; a *catalytic* bit vector of length B ; and one scratch area of $O(\log n)$ bits whose contents never survive a recursive call. Everything else is a suspended *local record* of a caller waiting for a lower query. The whole difficulty is to bound the sum of local records along the active call chain.

One cursor and paired ports. A product term $\mathcal{F}(x, z)\mathcal{B}(z, y)$ is evaluated with the cursor at x : a correction port moves it to z and returns a bounded inverse port, the right factor is queried at z with a task description that does not mention x , and the inverse port moves the cursor back. No second vertex address is ever stored across a recursive call (Theorem 11.1).

Incoming edges by tree contours. Some operations need the incoming entries of a column (the incoming gate of corrections, incoming ranks). These are found by viewing the deterministic program that maps a row port to its endpoint as a graph on program states. The component of a given success state is a tree directed toward it, whose leaves of the right kind are exactly the ports ending there. The Cook–McKenzie rotation/edge-flip contour [6] enumerates them without storing an address, as in reversible simulation of space-bounded computation [8] (Lemmas 11.5–11.7).

Catalytic key equality. At small slack, two endpoints are compared by their keys, but even a $\log B$ -bit key cannot be kept in every frame. The shared vector $V \in \{0, 1\}^B$ is used instead: visit v and flip $V_{k(v)}$; visit w and xor $V_{k(w)}$ into a result bit; visit v and flip $V_{k(v)}$ back; visit w and xor $V_{k(w)}$ into the result again. The result changes by exactly $\mathbf{1}[k(v) = k(w)]$ and V is restored, whatever its contents were. Only a phase and one bit are kept locally (Lemma 9.9). This is the transparent-program trick of catalytic computing [3].

Digit-by-digit arithmetic. Values are never stored in full. A query asks for one digit, in a redundant signed-digit stream with digits in $\{-6, \dots, 6\}$, obtained from integers N_t with $|N_t - 2^t z| < 2$ by $z_t = N_t - 2N_{t-1}$ (Lemma 10.1). Products, reciprocals (for the normalizations $1/c(q_y)$), ranks and sums are computed by circuits evaluated depth-first from regenerated gate addresses. A query for digit t that requests an operand digit j keeps at most $ad + b(t - j)$ local bits, with $j \leq t + cd$ (Theorem 10.9). So the numerical “clock” also telescopes.

4.2 The potential argument

Give every operation an *allowance*

$$w = K_1(1 + M) - (K_0 + s_0)d,$$

where M is its budget, d its slack and s_0 its position among a bounded number of operation strata. Every permitted descent (to a later stratum, to a coarser slack group, or to a child of budget at most $M - d$) lowers w by at least a constant times d . Combined with the numerical clock t , the potential

$$\Phi(w, t) = Aw + Dt$$

pays for each suspended segment: a frame waiting at (w_i, t_i) for a query at (w_{i+1}, t_{i+1}) keeps at most $\Phi(w_i, t_i) - \Phi(w_{i+1}, t_{i+1})$ bits. Summing along the chain telescopes to $\Phi(w_{\text{top}}, t_{\text{top}}) = O(M_{\text{top}} + t_{\text{top}}) = O(\log n)$. A second potential $\Psi = t + a'w$ never increases, which bounds every clock and therefore every scratch computation by $O(\log n)$ (proof of Theorem 11.1).

Why this beats $\log^{3/2} n$. In Savitch-style or Saks–Zhou-style recursion, every level of the recursion keeps a full vertex address or a full random seed. Here a level at slack d keeps $O(d)$ bits, and the levels’ slacks add up to the top budget. The cursor carries the only vertex address; the catalytic vector carries the only key table; the environment carries all the randomness, once. Each of these devices removes one of the $\Theta(\log n)$ -per-level costs.

Theorem 4.1 (Realization; Theorem 11.1 of [1]). *For every environment and every reachable task, every bit test and every digit query terminates, returns its defined value, restores the cursor and the catalytic vector, and uses $O(\log n)$ total space when the top budget and the requested clock are $O(\log n)$.*

5 Finishing the proof

5.1 From an average to one start, and the median

Lemma 5.1 (One start; Lemma 12.2 of [1]). *If $\frac{1}{N}\mathbb{E} \sum_v \frac{X_f(v)}{f} |\bar{W}(v) - W(v)|^8 \leq C_b^8 2^{-8p}$, then for any fixed z ,*

$$\Pr[|\bar{W}(z) - W(z)| > \eta \mid X_f(z) = 1] \leq N \frac{f}{\pi_f} C_b^8 2^{-8p} \eta^{-8}.$$

Proof. Keep only the summand $v = z$ and apply Markov's inequality. $\square$

Proof of Theorem 1.2 for fixed d (Theorem 12.1 of [1]). Let $r = d \lceil \log_2(n+2) \rceil$ and $\Delta = 2^{-r}$. Choose $L = O(\log n)$ so that $0 \leq p_{\mathcal{M}}(x) - W_L(z) \leq \Delta/8$ for the all-primary copy z of the start (Proposition 2.6). Choose the top budget $M_{\text{top}} = O(\log n)$ so large that $8p \geq 8r + C_N \log n + O(1)$. Then Lemma 5.1 with $\eta = \Delta/8$ shows that $|\bar{W}_L(z) - W_L(z)| \leq \Delta/8$ for at least $7/8$ of the environments that retain z .

Enumerate all environment encodings (field matrices with a full-rank frame, and raw salt strings) that retain z ; there are $2^{O(\log n)}$, each of $O(\log n)$ bits. For each, compute $K = r + 6$ signed digits of $\bar{W}_L(z)$ (Theorem 4.1), form the dyadic number v_σ , and clip it to $[0, 1]$; the truncation error is below $\Delta/8$. More than $3/4$ of the v_σ lie within $3\Delta/8$ of $p_{\mathcal{M}}(x)$, so their median does too. Find the median by binary search over $[0, 2^K]$, recounting with a fresh pass each time; only counters and the current environment are stored. Each digit query is a halting $O(\log n)$ -space computation, hence takes $2^{O(\log n)}$ steps, and there are $2^{O(\log n)}$ queries, so the time is polynomial. $\square$

Proof of Theorem 1.1. Apply the theorem with $d = 3$ to a BPL machine: the output is within $1/8$ of the acceptance probability, so comparing it with $1/2$ decides the language in L. Trivially $L \subseteq \text{RL} \subseteq \text{BPL}$. $\square$

5.2 Consequences

- **Promise problems.** The same algorithm is total, so $\text{PromiseBPL} = \text{PromiseL}$ (Cor. 12.3).
- **Accuracy as input.** Padding the input to length about $n2^q$ and running the fixed-accuracy algorithm on a virtual padded string gives error 2^{-q} in space $O(\log n + q)$ and time $n^{O(1)}2^{O(q)}$.
- **Finding accepting runs.** If $p_{\mathcal{M}}(x) \geq n^{-c}$, a deterministic log-space transducer outputs an accepting computation: estimate the continuation probability of every successor configuration to accuracy $n^{-c}/(8T(n))$ and follow a maximum (Cor. 12.4). The error per step is additive and deterministic, so no simultaneous-success event is needed.
- **An effective compiler.** One fixed randomized interpreter with enforced resource caps, together with the separator for it, yields an algorithm that, given a randomized machine

and claimed bounds $(n + 2)^a$ time and $b \log n$ space, outputs a deterministic machine with explicit bounds $K \log n$ space and $H(n + 2)^c$ time (Theorem 13.1).

5.3 The order of constants

As with every proof of this kind, correctness depends on choosing constants in the right order (Section 12.1 of the manuscript):

1. channel shape H, h_0 and the moment $s = 8$ (the walk gap is uniform in all later sizes);
2. correction thresholds, the cutoff Q , and the iteration count k_* ;
3. copy count D_0 , routing and gate thresholds, row-gate rank K ;
4. detector rank J , overload bound R_* , incoming rank K_c ;
5. precision coefficients, compression floors and capacities, walk length;
6. crude propagation constants, then block constants D_2, D_1 and the block width ε ;
7. controller allowance and guard constants;
8. the top-budget coefficient, the field-size exponent, and the minimum size of B .

6 A guide for the careful reader

6.1 What this note does and does not establish

This note describes the manuscript’s proof; it is not a verification. I checked the algebra of the hierarchy in [Section 2](#) (the transfer identity, the reward identity, the row-sum bound, the factorial truncation argument, the active-vertex count) and the final enumeration and median argument, and found them correct. For [Sections 3](#) and [4](#) I read the manuscript’s statements and arguments and have tried to report them faithfully, but those sections are long and intricate, and I did not check them line by line. There is no Lean formalization of this result in the repository.

6.2 Where to look hardest

1. **The hierarchy** (Section 4). It is elementary and the cleanest part. Its role is to reduce the problem to an exact object with $O(\log n)$ stages and bounded local structure. It is worth checking independently, since if it is right it is already a new structural fact about substochastic forward matrices.
2. **Shared randomness without independence** (Sections 5–7). Every task at every level uses the same $O(\log n)$ -bit environment. The proofs claim never to need independence between the errors of different tasks. The places to test this are products (Lemma 7.3), where a left factor’s support is used to sum a right factor’s error, and the use of conditional means on slices.
3. **Conditional averaging** (Section 6). Check that the validated walk variants really average exactly over the conditioned slice and that each sample function depends only on coordinates the walk varies (the “determinism hypothesis” of Lemma 7.8 and Lemma 9.2).

4. **Compression** (Section 8). The conditional-rarity argument (Lemma 8.4) applies to deterministic mean errors but not to noise; the proof handles noise separately using uniqueness of the first overflow. This split is delicate.
5. **Adaptive fingerprints** (Section 9). The salt-visibility rule must hold for every access operation, including inverse queries and port orders, not only for numerical values (Lemma 11.11 implements it). An error here would break the independence of the envelope from the block’s salt.
6. **Space accounting** (Sections 10–11). The claim is that every retained bit is charged to a decrease in allowance or clock. Look at the contour traversal of program graphs that include “opaque” lower queries, and at the guard construction for synthetic predecessor states.

6.3 Correspondence with the manuscript

This note	Manuscript
Theorems 1.1 and 1.2	Theorems 1.1, 1.2, 12.1
Lemma 2.1	Lemma 2.1 (Section 2)
Rank channels (Section 2.3)	Section 3, Lemma 3.1
Lemmas 2.2 to 2.4	Lemmas 4.1, 4.3, 4.4
Propositions 2.5 and 2.6	Propositions 4.5, 4.7
Atoms (Section 3.2)	Section 5
Averaging (Section 3.3)	Section 6
Budget schedule (Section 3.4)	Section 7
Compression (Section 3.5)	Section 8
Fingerprints (Section 3.6), Theorem 3.3	Section 9, Theorem 9.8
Digit streams	Section 10, Theorem 10.9
Theorem 4.1	Section 11, Theorem 11.1
Lemma 5.1 , finish, consequences	Section 12
Compiler	Section 13

References

- [1] OpenAI. *Exact Derandomization of Logarithmic Space: $L = RL = BPL$* . Manuscript, September 23, 2026. <https://github.com/openai/math>.
- [2] R. Aleliunas, R. Karp, R. Lipton, L. Lovász, C. Rackoff. Random walks, universal traversal sequences, and the complexity of maze problems. FOCS 1979.
- [3] H. Buhrman, R. Cleve, M. Koucký, B. Loff, F. Speelman. Computing with a full memory: catalytic space. STOC 2014.
- [4] J. L. Carter, M. N. Wegman. Universal classes of hash functions. *J. Comput. Syst. Sci.* 18(2), 1979.
- [5] K. Cheng, W. Hoza. Hitting sets give two-sided derandomization of small space. *Theory of Computing* 18, 2022.

- [6] S. Cook, P. McKenzie. Problems complete for deterministic logarithmic space. *J. Algorithms* 8(3), 1987.
- [7] W. Hoza. Better pseudodistributions and derandomization for space-bounded computation. *RANDOM* 2021.
- [8] K.-J. Lange, P. McKenzie, A. Tapp. Reversible space equals deterministic space. *J. Comput. Syst. Sci.* 60(2), 2000.
- [9] N. Nisan. Pseudorandom generators for space-bounded computation. *Combinatorica* 12(4), 1992.
- [10] N. Nisan. $RL \subseteq SC$. *Computational Complexity* 4, 1994.
- [11] O. Reingold. Undirected connectivity in log-space. *J. ACM* 55(4), 2008.
- [12] M. Saks, S. Zhou. $BP_HSPACE(S) \subseteq DSPACE(S^{3/2})$. *J. Comput. Syst. Sci.* 58(2), 1999.
- [13] Y. Shalom. Bounded generation and Kazhdan's property (T). *Publ. Math. IHÉS* 90, 1999.